
Same Top Fraction, Different Workload: Auditing Eligibility and Capacity in Clinical Risk Evaluation

Vache Oganisyan¹ Dmitry Lvov^{2,1} Ilya Pershin^{3,1}

¹Research Center of the Artificial Intelligence Institute,
Innopolis University, Innopolis, Russia

²ITMO University, St. Petersburg, Russia

³Kazan Federal University, Kazan, Russia

v.oganisian@innopolis.university d.lvov@innopolis.ru i.pershin@innopolis.ru

Abstract

Clinical risk scores turn patient rankings into finite worklists for specialist review, monitoring, or prevention. Yet cross-cohort evaluations often compare the same top fraction across candidate sets of different sizes, changing both who can be ranked and how many people receive service. We introduce a decision-aware four-cell audit that keeps the full-population event denominator fixed and crosses two eligibility rules with two absolute budgets, separating broader reach from additional capacity. We apply it to frozen birth-record score vectors for preterm birth, neonatal intensive care admission, and low birth weight in approximately 3.48 million CDC Natality 2023 records. At the nominal top 10% operating point, moving from early-entry eligibility to a retrospective all-record upper bound reports a 39–42% relative increase in captured events while expanding the worklist by one third. Symmetric accounting allocates about three fifths of the joint gain to added capacity and two fifths to broader eligibility; at a fixed budget, the eligibility gain is less than half the cohort-relative estimate. In 22% of controlled conditions, the cohort-relative comparison improves while fixed-budget capture declines. The audit therefore distinguishes population reach from resource expansion and makes the policy represented by a top-fraction metric explicit. Code and reproduction materials will be publicly released at <https://github.com/vache4ogan/prenatal-risk-scoreability>.

1 Introduction

Clinical risk models increasingly determine which records receive specialist review, monitoring, outreach, or preventive services [Annarumma et al., 2019, Baltruschat et al., 2021, Jung et al., 2021, Singh et al., 2023]. These workflows are capacity constrained: a model produces a ranking, but an organization acts on a finite number of cases. Top- q precision, recall, or event capture therefore appears to connect predictive performance directly to an operational decision. Prenatal case-management studies, for example, evaluate the highest-risk 5–15% because program capacity covers a fraction of eligible pregnancies [Pan et al., 2017].

The concrete scale of this policy difference is visible in CDC Natality 2023 [National Center for Health Statistics, 2026]. An early-entry candidate set of approximately 2.60 million records yields a top-10% worklist of about 260,000, whereas an all-record set of approximately 3.48 million yields about 348,000 selections at the same nominal fraction. The comparison therefore changes both the candidate gate and the absolute workload. Clinical decision frameworks distinguish predictive




Figure 1: The same nominal top 10% encodes a different workload. Panel (a) summarizes the target-specific CDC 2023 cohorts: moving from the early-entry gate to the all-record upper bound increases the selected count from about 260,000 to 348,000. Panel (b) allocates the associated increase in population event capture at $q = 10\%$. Values inside bars are percentage-point allocations; capacity contributes about 62% of each total gain.

ranking from action under finite resources [Jung et al., 2021, Singh et al., 2023]: a fixed-capacity service and a service that expands with candidate volume implement different policies and require different contrasts.

Prenatal risk triage makes this ambiguity visible. Pregnancies enter care at different times, and adverse outcomes are concentrated among records with no documented prenatal care or unknown start timing. An early-entry score ranks events inside its candidate gate; events outside the gate contribute zero reachable cases regardless of discrimination. National vital statistics retain both the eligible and excluded outcomes, making it possible to distinguish retrieval within the gate from reach across the complete target population.

We audit preterm birth, neonatal intensive care unit (NICU) admission, and low birth weight in CDC Natality. Figure 1 previews how the workload change alters the interpretation of the reported gain.

Our contributions are:

1. **An auditable decision estimand.** A shared score vector and population event denominator connect a same-fraction claim to four explicit policies. The two fixed-budget contrasts reveal which part survives equal workload; their interaction measures dependence on the chosen budget.
2. **A national-scale measurement study.** Frozen 2022 scoring pipelines applied to CDC 2023 show that the same top 10% expands workload by 33.8% and that capacity receives 61.6–62.2% of the symmetric allocation across three outcomes.
3. **A mechanism map and reporting profile.** An 81-condition stress test locates sign-reversal regimes, and the resulting checklist makes the candidate gate, absolute workload, and event denominator explicit.

2 Related Work

Risk ranking under limited resources. Prenatal case-management, imaging worklists, and clinical workflow studies connect risk scores to finite review or treatment capacity [Pan et al., 2017, Annarumma et al., 2019, Baltruschat et al., 2021, Jung et al., 2021, Singh et al., 2023]. These studies connect operating points to the capacity and value of a clinical workflow. Resource-constrained net benefit evaluates whether the expected value of acting on a score justifies its workload [Singh et al., 2023]. Calibrated subset selection constructs shortlists containing a desired number of qualified candidates in expectation [Wang et al., 2022]. These objectives govern selection within a specified population and resource model. Interpreting a reported gain across candidate pools additionally requires tracking the absolute budget induced by each pool’s size.

Selective access to clinical data and services. Clinical records reflect who enters care, is tested, and receives a documented label. Informed presence and selection bias can alter both apparent risk and the population represented by an EHR cohort [Haneuse and Daniels, 2016, Goldstein et al., 2016, Sisk et al., 2021]; disparate censorship describes related variation in testing and label availability [Chang et al., 2022]. Health-service coverage frameworks separate service availability from realized coverage [Tanahashi, 1978]. This literature establishes that the candidate gate is informative, but a top-fraction comparison adds a second moving quantity because the number served scales with the observed cohort.

Evaluation validity and decomposition. Clinical reporting guidance requires explicit target populations, intended use, discrimination, and calibration [Collins et al., 2024]. Symmetric allocation across factor orderings is established decomposition machinery [Kitagawa, 1955, Shapley, 1953, Shorrocks, 2013]. The decomposition rule specifies how a joint difference is allocated once the comparison cells have been defined. Cross-cohort top-fraction evaluation still requires a choice of those cells: whose outcomes form the denominator, which records can enter the ranking, and how many positions each policy funds. Without these choices, a constant percentage can conceal a changing workload and leave the fixed-capacity conclusion unresolved.

3 Methodology

The audit fixes the target population, score vector, and event denominator, then crosses two candidate-set rules with two absolute budgets. Figure 2 shows the resulting four cells and the two paths whose average yields the symmetric allocation.




Figure 2: Eligibility–capacity audit. Rows vary the candidate set and columns vary the absolute selection budget while the score vector and full-population event denominator remain fixed. Blue arrows change candidate eligibility; orange arrows change capacity. Averaging the two paths from C_{00} to C_{11} gives the symmetric Shapley allocation, whereas $C_{11} - C_{00}$ is the usual same- q comparison and changes both factors.

3.1 Target population and scoring study

We use CDC Natality public-use live-birth records and official field definitions [National Center for Health Statistics, 2026, 2017]. CDC 2022 supplies model development with an 80/20 train/validation split; the frozen pipelines generate the shared score vectors analyzed in CDC 2023. Model fitting and hyperparameter selection use CDC 2022 exclusively; CDC 2023 is the temporal cohort in which the audit hypothesis was developed and evaluated. The target population contains U.S.-resident singleton live births with a known outcome: preterm birth below 37 completed obstetric-estimate weeks ($N = 3,480,382$, 303,001 events), NICU admission ($N = 3,476,584$, 303,285 events), and birth weight below 2,500 g ($N = 3,480,316$, 243,776 events).

Prenatal-care month defines four mutually exclusive groups: early entry in months 1–3, later documented entry in months 4–10, no prenatal care, and unknown start timing. The early candidate set contains months 1–3. The documented-start set contains months 1–10. The all-record set adds no-care and unknown-timing records and therefore bounds candidate and event availability under retrospective complete scoring. The care-month rule operationalizes documented entry into prenatal

care rather than feature-level EHR availability; the birth-certificate model serves as a shared retrospective ranking instrument.

The baseline-oriented score uses nine maternal birth-certificate variables: race, height, pre-pregnancy BMI and weight, previous live births now living, previous live births now dead, other pregnancy outcomes, pre-pregnancy diabetes, and pre-pregnancy hypertension. Outcome, delivery, and neonatal fields are excluded. Numeric variables use development-set median imputation and standardization; race uses most-frequent imputation and one-hot encoding. One LightGBM model per outcome uses 350 trees, learning rate 0.05, 31 leaves, minimum child size 100, L2 penalty 1, and seed 2026 [Ke et al., 2017]. The audit compares candidate and budget policies on a shared score vector. A camera-ready replication applies the same frozen 2022 pipelines to CDC 2024, with identical cohort rules, four cells, and 500 paired resamples per outcome.

3.2 Decision metrics

Let $\mathcal{P}$ contain N records and $E = \sum_{i \in \mathcal{P}} Y_i$ events. For candidate set S , event count E_S , and integer budget b , deterministic tie breaking defines the b highest scores $\text{Top}(S, b)$. Their true-positive count is

$$\text{TP}(S, b) = \sum_{i \in \text{Top}(S, b)} Y_i. \quad (1)$$

We report both within-gate performance and population reach:

$$\text{Precision}(S, b) = \frac{\text{TP}(S, b)}{b}, \quad \text{Recall}_{\text{eligible}}(S, b) = \frac{\text{TP}(S, b)}{E_S}, \quad (2)$$

$$\text{Capture}(S, b) = \frac{\text{TP}(S, b)}{E}, \quad \text{Ceiling}(S) = \frac{E_S}{E}. \quad (3)$$

Population capture factors into retrieval and reach,

$$\text{Capture}(S, b) = \text{Recall}_{\text{eligible}}(S, b) \text{ Ceiling}(S). \quad (4)$$

Each result therefore reports target-population N and E , eligible n and E_S , exact selected n , precision, eligible recall, population capture, and the event ceiling. ROC-AUC and average precision (AP) provide a ranking-signal check; the four-cell contrasts carry the primary decision claim.

3.3 Four-cell eligibility–capacity audit

Let $S_e \subseteq S_a$ be early and all-record candidate sets. A nominal fraction q induces $b_e = \text{round}(q|S_e|)$ and $b_a = \text{round}(q|S_a|)$. With $b_a \leq |S_e|$, the complete design is

$$C_{uv} = \text{Capture}(S_u, b_v), \quad u, v \in \{0, 1\}, \quad (5)$$

where 0 denotes early and 1 denotes all-record eligibility or capacity. The usual cohort-relative contrast is $\Delta = C_{11} - C_{00}$. Its interaction is $I = C_{11} - C_{10} - C_{01} + C_{00}$. We retain both fixed-budget eligibility contrasts and summarize the joint change with the standard symmetric Shapley allocation [Shapley, 1953]:

$$\begin{aligned} \phi_E &= \frac{1}{2} [(C_{10} - C_{00}) + (C_{11} - C_{01})], \\ \phi_B &= \frac{1}{2} [(C_{01} - C_{00}) + (C_{11} - C_{10})]. \end{aligned} \quad (6)$$

The exact identity $\Delta = \phi_E + \phi_B$ holds, and $s_B = \phi_B/\Delta$ gives the symmetric capacity share for positive Δ . For a frozen deterministic ranking, increasing b adds nested positions, so the capacity component is non-negative; expanding S at fixed b can have either sign. A uniformly random ranking supplies the mechanical reference $\mathbb{E}[\text{Capture}(S, b)] = (b/|S|)(E_S/E)$.

The contribution is the comparison design: off-diagonal cells test equal-workload claims against a denominator retaining events outside the gate. The established Shapley rule summarizes these event-count changes; clinical effects require a separate intervention model.

3.4 Evaluation and mechanism stress test

The empirical audit uses $q \in \{0.05, 0.10\}$. Point estimates are accompanied by 500 paired full-size empirical resamples per outcome and q . Each resample uses common record multiplicities across all four cells while holding the observed scores and absolute budgets fixed. The resulting central 95% ranges measure conditional stability of the reported four-cell contrasts.

The simulation draws latent risk $Z \sim \mathcal{N}(0, 1)$, outcomes with $P(Y = 1 \mid Z) = \text{logit}^{-1}(\alpha_Y + Z)$, eligibility with $P(A = 1 \mid Z) = \text{logit}^{-1}(\alpha_A + \gamma Z)$, and scores $R = Z + \sigma\epsilon$, where $\epsilon \sim \mathcal{N}(0, 1)$ is independent. Conditional on Z , outcome and eligibility draws are independent. Intercepts set marginal event prevalence to 0.10 and availability to a . Crossing $a \in \{0.5, 0.7, 0.9\}$, $\gamma \in \{-1, 0, 1\}$, $\sigma \in \{0.5, 1, 2\}$, and $q \in \{0.05, 0.10, 0.20\}$ gives 81 conditions. Each of 100 independent repetitions contains 100,000 evaluation records and 50,000 independent threshold-selection records, with shared random draws across conditions. Appendix C specifies all parameters, coupling, and generation steps.

4 Results

4.1 Candidate eligibility and event ceilings

Early entry contains about 74.7% of target-population records, but only 71.8% of preterm events, 70.6% of NICU events, and 69.4% of low-birth-weight events. Before ranking begins, the early gate therefore leaves 28.2–30.6% of events outside its reachable population. Including every documented prenatal-care start in months 1–10 raises the event ceiling to 91.1–92.1%.

The excluded records carry concentrated risk. Outcome prevalence among no-care records is 2.4–2.7 times the early-entry prevalence across the three targets; unknown-timing records are 1.7–1.9 times higher. Candidate expansion therefore changes both population reach and case mix, making a full-population event denominator essential.

Table 1: The candidate gate fixes reach before ranking. Early entry denotes prenatal-care initiation in months 1–3; documented start includes months 1–10. Risk ratios compare outcome prevalence in excluded groups with the early-entry group.

Outcome	Early eligible n	Early share	Early ceiling	Documented ceiling	No-care / unknown risk
Preterm birth	2,601,060	74.7%	71.8%	91.1%	$2.7\times / 1.8\times$
NICU admission	2,597,274	74.7%	70.6%	92.1%	$2.4\times / 1.7\times$
Low birth weight	2,600,093	74.7%	69.4%	91.6%	$2.6\times / 1.8\times$

4.2 Eligibility–capacity accounting under top- q selection

Figure 1 and Table 2 connect the nominal operating point to its actual workload. Across outcomes, the selected count increases from approximately 260,000 to 348,000. Population event capture rises from 14.63% to 20.30% for preterm birth, 13.93% to 19.52% for NICU, and 16.06% to 22.83% for low birth weight. These are relative gains of 38.8%, 40.1%, and 42.2%, respectively.

The proportional policy adds 87,931–88,023 service slots per outcome and captures 16,505–17,177 additional events. The yield of this combined policy change is 18.8–19.5 additional events per 100 added slots, equivalent to 5.1–5.3 additional selections per extra captured event. This ratio includes both expanded eligibility and the 33.8% larger capacity. Capacity-only yields compare columns within a fixed row of the four-cell design.

The four-cell audit changes the interpretation of those gains. Symmetric accounting assigns 2.11–2.60 percentage points to eligibility and 3.47–4.17 points to additional capacity, placing 61.6–62.2% of the joint contrast on the workload increase. Holding the early budget fixed yields eligibility gains of 1.89–2.27 points; holding the all-record budget fixed yields 2.32–2.93 points. Both are materially smaller than the joint 5.59–6.77-point cohort-relative contrast.

The pattern persists at $q = 5\%$, where capacity receives 63.9–65.6% of the symmetric allocation. Conditional ranges remain above 60.4% for every target and operating point. A random ranking

Table 2: Same nominal top 10%, different operational result. Population capture uses the full target-population event denominator. ϕ_E and ϕ_B are the symmetric eligibility and capacity allocations; conditional empirical-resampling ranges appear in Appendix 7.

Outcome	Selected n	Workload	Capture (%)	Rel. gain	ϕ_E/ϕ_B (pp)	s_B
Preterm birth	260,106→348,038	+33.8%	14.63→20.30%	+38.8%	2.15/ 3.52	62.1%
NICU admission	259,727→347,658	+33.9%	13.93→19.52%	+40.1%	2.11/ 3.47	62.2%
Low birth weight	260,009→348,032	+33.9%	16.06→22.83%	+42.2%	2.60/ 4.17	61.6%

allocates 79.7–87.9% to capacity at $q = 10\%$, showing that a majority capacity allocation can arise mechanically under a proportional budget. The learned ranking shifts more of the joint gain toward eligibility because the added candidate groups contain greater event mass. Complete four-cell counts and both operating points appear in Appendix B.

4.3 Fixed-budget and transported-threshold policies

Table 3 compares three policies on the same frozen score vector. Cohort-relative top-10% evaluation jointly changes eligibility and workload. Fixing the all-record absolute budget isolates candidate expansion: capture increases by 2.32–2.93 points, with exactly the same number selected on both sides. A threshold chosen once on the CDC 2022 validation split selects 9.15–9.82% of early-entry records and 10.20–10.40% of all records in CDC 2023. Under that transported threshold, the all-record worklist is 1.42–1.49 times larger and capture increases by 6.36–8.19 points.

Table 3: One score vector, three resource policies at the nominal 10% operating point. Counts and population capture compare early-entry with all-record eligibility. The score threshold is chosen on CDC 2022 and transported unchanged to CDC 2023.

Outcome	Resource policy	Selected n : early → all	Capture: early → all	Δ (pp)
Preterm birth	Cohort-relative top 10%	260,106→348,038	14.63→20.30%	+5.67
	Fixed absolute b_a	348,038→348,038	17.89→20.30%	+2.40
	Transported 2022 threshold	245,487→361,709	14.06→20.84%	+6.78
NICU admission	Cohort-relative top 10%	259,727→347,658	13.93→19.52%	+5.59
	Fixed absolute b_a	347,658→347,658	17.20→19.52%	+2.32
	Transported 2022 threshold	255,166→361,714	13.76→20.12%	+6.36
Low birth weight	Cohort-relative top 10%	260,009→348,032	16.06→22.83%	+6.77
	Fixed absolute b_a	348,032→348,032	19.90→22.83%	+2.93
	Transported 2022 threshold	237,973→355,034	14.98→23.17%	+8.19

Transporting the full-precision 2022 cutoffs to CDC 2024 selects 10.60% of preterm, 10.64% of NICU, and 10.20% of low-birth-weight records: 21,112, 22,400, and 6,952 more slots than a 10% budget. The complete 2024 four-cell replication assigns 60.9–62.8% of the top-10% joint gain to capacity, with conditional 95% ranges spanning 59.8–63.9% across outcomes (Appendix B). Both fixed-budget eligibility contrasts remain positive for every outcome and q . This replication uses the original numeric race encoding and supersedes the reviewed 2024 threshold-only calculation, whose string encoding silently mapped race categories to unknown values.

4.4 Controlled mechanism stress test

Figure 3 maps the protocol difference across all 81 synthetic conditions. Eighteen conditions report a positive cohort-relative gain while candidate expansion at the all-record absolute budget reduces population capture. Nine have a central replicate range entirely below zero. Every reversal arises when higher-risk records are preferentially available, so the newly admitted pool is lower risk and can displace selected events.

At 70% availability, medium score quality, and $q = 10\%$, for example, cohort-relative capture increases 5.52 points while fixed-budget candidate expansion changes capture by -0.28 points; 93 of 100 repetitions reverse. The largest mean absolute interaction is 9.49 points. In CDC 2023, both fixed-budget eligibility contrasts remain positive for every outcome and q , so the ambiguity




Figure 3: Full-grid mechanism stress test. Panel (a) compares the cohort-relative conclusion with candidate expansion at the all-record-derived absolute budget b_a . The shaded quadrant contains sign reversals: the joint top- q contrast is positive while fixed-budget eligibility expansion is negative. Panel (b) gives replicate-level reversal frequency for conditions where higher risk is more available. Marker shape denotes score quality.

changes effect magnitude. The controlled grid shows that selective availability can also change effect direction.

4.5 Temporal score diagnostics

The frozen 2022 pipelines retain similar discrimination in CDC 2023 (Table 4). Across outcomes, ROC-AUC and average precision change by at most 0.003 from 2022 validation to 2023. Brier skill relative to a prevalence-only forecast is 1.7–2.4%. The four-cell audit fixes each transported score vector, isolating within-year policy contrasts from temporal score drift.

Table 4: Frozen-score temporal check. Brier baseline is the constant CDC 2023 prevalence prediction; skill is $1 - \text{Brier}_{\text{model}} / \text{Brier}_{\text{baseline}}$.

Outcome	ROC-AUC: 2022 → 2023	AP: 2022 → 2023	Brier: model / baseline	Brier skill
Preterm birth	0.614→0.612	0.146→0.146	0.0780 / 0.0795	1.9%
NICU admission	0.611→0.608	0.139→0.141	0.0783 / 0.0796	1.7%
Low birth weight	0.653→0.652	0.133→0.132	0.0636 / 0.0651	2.4%

4.6 Sensitivity and robustness

Two feature ablations preserve the allocation pattern. Removing maternal race and refitting each outcome model at $q = 10\%$ leaves capacity shares of 60.9–62.4%. A separate eight-feature specification removes prior other pregnancy outcomes (prior_terminations) and repeats the complete audit at both budgets with paired resampling. Capacity still accounts for 61.5–64.3% of the joint contrast, and all interaction ranges remain above zero (+0.32 to +0.69 percentage points). The result therefore persists across two alternative conservative landmark specifications.

5 Discussion

Resource interpretation of top- q evaluation. Top- q evaluation is coherent when service volume scales with the candidate population; a fixed budget answers the constrained-capacity question. Here, a 38.8–42.2% capture gain includes 33.8% more workload. Making this policy visible helps prevent unplanned demand from redirecting scarce services.

From audit to policy choice. For a fixed service budget B , compare feasible policies at the same selected count and maximize captured events when case values and selection costs are equal. With

flexible capacity, an explicit screening utility connects the audit to resource-constrained decision analysis [Singh et al., 2023]:

$$U(S, b) = v \text{ TP}(S, b) - cb - K(S), \quad (7)$$

where $v > 0$ values a captured event, c costs one selection, and $K(S)$ is the cost of making the candidate pool reachable. Prefer a feasible policy change when $v \Delta \text{ TP} > c \Delta b + \Delta K$. For the joint CDC change and $\Delta K = 0$, the break-even cost-to-value ratio is 0.188–0.195. These retrospective retrieval values require intervention benefits and harms to translate into prospective clinical utility.

Applicability across deterministic score vectors. Nested top- b selections imply capacity monotonicity for any fixed ranking. The numerical allocation depends on case mix, ranking quality, eligibility, and budget: 61.6–62.2% characterizes these three CDC 2023 tasks and LightGBM scores. Feature ablations test specification sensitivity; the controlled grid shows how availability mechanisms can reverse fixed-budget contrasts.

A recommended reporting profile. Table 5 specifies the minimum reporting items.

Table 5: Minimum reporting checklist for cross-cohort top- q evaluation. The final column illustrates each item with the CDC 2023 audit.

Required item	Minimum report	Example from this study
Target population	Population rule, N , total events E , and event denominator	CDC 2023 singleton births; $N \approx 3.48\text{M}$; all target events remain in the denominator
Candidate gate	Eligibility rule, eligible n , E_S , and event ceiling E_S/E	Care entry in months 1–3; $n \approx 2.60\text{M}$; ceiling 69.4–71.8%
Resource rule	Fraction q or absolute budget b , with exact selected n	Top 10%; selected n : 260k→348k (+33.8% workload)
Operating result	Precision, eligible recall, and population event capture	Capture: 13.93–16.06%→19.52–22.83%
Cross-cohort audit	Four cells, fixed-budget contrasts, interaction, and symmetric allocation	Capacity receives 61.6–62.2% of the joint allocation

Implementation in existing evaluation pipelines. Stored scores, outcomes, and candidate masks construct the cells; common resampling weights preserve pairing without releasing individual records.

Scope and next validation. Birth-certificate records provide retrospective outcomes and a care-entry eligibility proxy. Prospective EHR landmarks would establish when each input and intervention becomes available. The CDC 2024 replication supports temporal stability within this source; additional model families and independent clinical settings would test broader transport. Policy evaluation with heterogeneous benefits, acquisition costs, or subgroup priorities can use the same four-cell counts under an explicitly chosen utility.

6 Conclusion

Top-fraction comparisons encode service decisions as well as ranking quality. In CDC 2023, the same top 10% expands workload by one third, with about 62% of the capture gain allocated to capacity; the CDC 2024 replication preserves this pattern. Controlled experiments also identify opposite-sign conclusions. Reporting the population, candidate rule, selected count, event denominator, and fixed-budget contrast makes these policy claims auditable.

Acknowledgments and Disclosure of Funding

The study was supported by the Ministry of Economic Development of the Russian Federation (agreement No. 139-10-2025-034 dd. 19.06.2025, IGG 000000C313925P4D0002).

References

- Mauro Annarumma, Samuel J. Withey, Robert J. Bakewell, Emanuele Pesce, Vicky Goh, and Giovanni Montana. Automated triaging of adult chest radiographs with deep artificial neural networks. *Radiology*, 291(1):196–202, 2019. doi: 10.1148/radiol.2018180921.
- Ivo M. Baltruschat, Leonhard Steinmeister, Hannes Nickisch, Axel Saalbach, Michael Grass, Gerhard Adam, Tobias Knopp, and Harald Ittrich. Smart chest X-ray worklist prioritization using artificial intelligence: A clinical workflow simulation. *European Radiology*, 31(6):3837–3845, 2021. doi: 10.1007/s00330-020-07480-7.
- Trenton Chang, Michael W. Sjoding, and Jenna Wiens. Disparate censorship and undertesting: A source of label bias in clinical machine learning. In *Proceedings of the 7th Machine Learning for Healthcare Conference*, volume 182 of *Proceedings of Machine Learning Research*, pages 343–390, 2022. URL <https://proceedings.mlr.press/v182/chang22a.html>.
- Gary S. Collins, Karel G. M. Moons, Paula Dhiman, Richard D. Riley, Andrew L. Beam, Ben Van Calster, Marzyeh Ghassemi, Xiaoxuan Liu, Johannes B. Reitsma, Maarten van Smeden, et al. TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods. *BMJ*, 385:e078378, 2024. doi: 10.1136/bmj-2023-078378.
- Benjamin A. Goldstein, Nrupen A. Bhavsar, Matthew Phelan, and Michael J. Pencina. Controlling for informed presence bias due to the number of health encounters in an electronic health record. *American Journal of Epidemiology*, 184(11):847–855, 2016. doi: 10.1093/aje/kww112.
- Sebastien Haneuse and Michael J. Daniels. A general framework for considering selection bias in EHR-based studies: What data are observed and why? *eGEMs*, 4(1):1203, 2016. doi: 10.13063/2327-9214.1203.
- Kenneth Jung, Sehj Kashyap, Anand Avati, Stephanie Harman, Heather Shaw, Ron C. Li, Margaret Smith, Kenny Shum, Jacob Javitz, Yohan Vetteth, Tina Seto, Steven C. Bagley, and Nigam H. Shah. A framework for making predictive models useful in practice. *Journal of the American Medical Informatics Association*, 28(6):1149–1158, 2021. doi: 10.1093/jamia/ocaa318.
- Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL https://proceedings.neurips.cc/paper_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html.
- Evelyn M. Kitagawa. Components of a difference between two rates. *Journal of the American Statistical Association*, 50(272):1168–1194, 1955. doi: 10.1080/01621459.1955.10501299.
- Microsoft Corporation and LightGBM Contributors. Lightgbm. GitHub repository, MIT License, 2026. URL <https://github.com/lightgbm-org/LightGBM>. Accessed 2026-08-26.
- National Center for Health Statistics. Facility worksheet for the live birth certificate. Centers for Disease Control and Prevention, 2017. URL <https://www.cdc.gov/nchs/data/dvs/facility-worksheet-2016.pdf>. Final 12/16 version; accessed 2026-08-26.
- National Center for Health Statistics. Nchs data user agreement. Centers for Disease Control and Prevention, 2025. URL <https://www.cdc.gov/nchs/policy/data-user-agreement.html>. Accessed 2026-08-26.
- National Center for Health Statistics. Vital statistics online data portal: Natality public-use data files and documentation. Centers for Disease Control and Prevention, 2026. URL https://www.cdc.gov/nchs/data_access/vitalstatsonline.htm. Accessed 2026-08-26.
- Ian Pan, Laura B. Nolan, Rashida R. Brown, Romana Khan, Paul van der Boor, Daniel G. Harris, and Rayid Ghani. Machine learning for social services: A study of prenatal case management in Illinois. *American Journal of Public Health*, 107(6):938–944, 2017. doi: 10.2105/AJPH.2017.303711.

- Lloyd S. Shapley. A value for n -person games. In Harold W. Kuhn and Albert W. Tucker, editors, *Contributions to the Theory of Games II*, pages 307–317. Princeton University Press, Princeton, 1953. doi: 10.1515/9781400829156-012.
- Anthony F. Shorrocks. Decomposition procedures for distributional analysis: A unified framework based on the shapley value. *The Journal of Economic Inequality*, 11(1):99–126, 2013. doi: 10.1007/s10888-011-9214-z.
- Karandeep Singh, Nigam H. Shah, and Andrew J. Vickers. Assessing the net benefit of machine learning models in the presence of resource constraints. *Journal of the American Medical Informatics Association*, 30(4):668–673, 2023. doi: 10.1093/jamia/ocad006.
- Rose Sisk, Lijing Lin, Matthew Sperrin, Jessica K. Barrett, Brian Tom, Karla Diaz-Ordaz, Niels Peek, and Glen P. Martin. Informative presence and observation in routine health data: A review of methodology for clinical risk prediction. *Journal of the American Medical Informatics Association*, 28(1):155–166, 2021. doi: 10.1093/jamia/ocaa242.
- T. Tanahashi. Health service coverage and its evaluation. *Bulletin of the World Health Organization*, 56(2):295–303, 1978. URL <https://pmc.ncbi.nlm.nih.gov/articles/PMC2395571/>.
- Lequn Wang, Thorsten Joachims, and Manuel Gomez Rodriguez. Improving screening processes via calibrated subset selection. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 22702–22726, 2022. URL <https://proceedings.mlr.press/v162/wang22j.html>.

A Outcome, Model, and Cohort Details

Table 6: Baseline-oriented birth-certificate variables used by the frozen ranking instrument. Encounter-specific availability at months 1–3 requires validation in a prospective EHR source.

Feature	CDC source meaning
Maternal race	Maternal characteristic on birth record
Maternal height	Maternal characteristic on birth record
Pre-pregnancy BMI	Derived from pre-pregnancy measures
Pre-pregnancy weight	Weight before pregnancy
Previous live births now living	Previous live-born infants still living
Previous live births now dead	Previous live-born infants no longer living
Other pregnancy outcomes	Non-live-birth outcomes; may include current pregnancy
Pre-pregnancy diabetes	Diabetes diagnosed before this pregnancy
Pre-pregnancy hypertension	Hypertension diagnosed before this pregnancy

For CDC 2023 preterm records, early entry contains 2,601,060 records and 217,632 events; later entry contains 739,298 and 58,282; no care contains 76,520 and 17,387; unknown timing contains 63,504 and 9,700. Target-specific counts differ slightly because each outcome has its own missing-label exclusions. Source code 99 is normalized to missing before cohort construction, while source value 0 denotes no prenatal care.

B Complete Four-Cell Results

The point-estimate budgets at $q = 5\%$ are 130,053/174,019 for preterm, 129,864/173,829 for NICU, and 130,005/174,016 for low birth weight (early/all-record). At $q = 10\%$ they are 260,106/348,038, 259,727/347,658, and 260,009/348,032. These observed integer budgets remain fixed across empirical resamples.

The camera-ready replication uses 3,512,780 preterm records (307,612 events), 3,510,341 NICU records (310,650 events), and 3,514,422 low-birth-weight records (245,967 events) from CDC 2024. Models, preprocessing, and thresholds remain frozen from 2022. The corrected numeric MRACE31 representation matches model fitting; replay on CDC 2023 reproduces all 24 original four-cell point estimates exactly. Both full-precision and six-decimal threshold results, all 24 replication cells, and 3,000 paired-resample rows are archived.

Table 7: Complete symmetric-allocation results. All effects are percentage-point changes in population event capture; brackets give central 95% paired empirical-resampling ranges conditional on frozen scores and observed budgets.

Outcome	q	Joint Δ	ϕ_E	ϕ_B	Symmetric share s_B	Interaction
Preterm birth	5%	3.37 [3.30, 3.43]	1.16 [1.09, 1.22]	2.21 [2.17, 2.25]	65.6 [64.3, 67.1]%	0.33 [0.26, 0.41]
Preterm birth	10%	5.67 [5.58, 5.77]	2.15 [2.06, 2.24]	3.52 [3.48, 3.57]	62.1 [61.1, 63.3]%	0.51 [0.44, 0.60]
NICU admission	5%	3.31 [3.26, 3.38]	1.19 [1.13, 1.25]	2.12 [2.09, 2.16]	64.1 [62.8, 65.4]%	0.29 [0.23, 0.36]
NICU admission	10%	5.59 [5.51, 5.67]	2.11 [2.03, 2.20]	3.47 [3.44, 3.52]	62.2 [61.1, 63.3]%	0.42 [0.34, 0.51]
Low birth weight	5%	3.96 [3.89, 4.06]	1.43 [1.34, 1.52]	2.53 [2.50, 2.57]	63.9 [62.3, 65.4]%	0.40 [0.32, 0.49]
Low birth weight	10%	6.77 [6.67, 6.89]	2.60 [2.49, 2.72]	4.17 [4.12, 4.22]	61.6 [60.4, 62.7]%	0.66 [0.55, 0.78]

Table 8: Auditable four-cell point estimates. Each cell reports population event capture in percent followed by selected true events in parentheses. The two budgets are exact selected counts.

Outcome	q	Budgets b_e/b_a	$C_{00}: S_e, b_e$	$C_{01}: S_e, b_a$	$C_{10}: S_a, b_e$	$C_{11}: S_a, b_a$
Preterm birth	5%	130,053/174,019	9.05 (27,434)	11.10 (33,618)	10.04 (30,432)	12.42 (37,631)
Preterm birth	10%	260,106/348,038	14.63 (44,326)	17.89 (54,216)	16.52 (50,061)	20.30 (61,503)
NICU admission	5%	129,864/173,829	8.51 (25,801)	10.48 (31,795)	9.55 (28,967)	11.82 (35,854)
NICU admission	10%	259,727/347,658	13.93 (42,252)	17.20 (52,160)	15.84 (48,033)	19.52 (59,201)
Low birth weight	5%	130,005/174,016	9.53 (23,229)	11.86 (28,918)	10.76 (26,224)	13.49 (32,892)
Low birth weight	10%	260,009/348,032	16.06 (39,145)	19.90 (48,505)	18.33 (44,679)	22.83 (55,650)

C Synthetic Grid Details

Data-generating process. Let $g(x) = (1 + e^{-x})^{-1}$. For each evaluation record draw mutually independent $Z_i, \epsilon_i \sim \mathcal{N}(0, 1)$ and $U_i^Y, U_i^A \sim \text{Uniform}(0, 1)$. Define

$$Y_i = \mathbf{1}\{U_i^Y < g(\alpha_Y + Z_i)\}, \quad R_i = Z_i + \sigma\epsilon_i, \quad (8)$$

$$A_i = \mathbf{1}\{U_i^A < g(\alpha_A(a, \gamma) + \gamma Z_i)\}. \quad (9)$$

The full population contains all records and the available set is $\{i : A_i = 1\}$. The intercepts solve $\mathbb{E}[g(\alpha_Y + Z)] = 0.10$ and $\mathbb{E}[g(\alpha_A + \gamma Z)] = a$ for $Z \sim \mathcal{N}(0, 1)$. The outcome-risk coefficient is 1.0; $\alpha_Y = -2.5642215001$. Availability slopes $\gamma = -1, 0, 1$ encode higher risk being less available, independent of availability, or more available, respectively.

The implementation uses 80-node Gauss–Hermite quadrature and bisection on $[-40, 40]$ with marginal-probability tolerance 10^{-13} and at most 200 iterations. Table 10 displays rounded values; the code solves them at runtime. The grid crosses $a \in \{0.5, 0.7, 0.9\}$, $\gamma \in \{-1, 0, 1\}$, $\sigma \in \{0.5, 1, 2\}$, and $q \in \{0.05, 0.10, 0.20\}$.

Repetitions and shared random draws. NumPy SeedSequence(2028).spawn(100) supplies independent repetition seeds to default_rng. Each repetition draws, in order, 100,000 evaluation risks, outcome uniforms, availability uniforms, and standard-normal score errors, followed by 50,000 independent threshold-selection risks and errors. The same base draws are reused across the entire grid: changing γ or a changes the availability rule, changing σ rescales the score error, and changing q changes the budget. Outcomes, scores, and masks remain shared between the four cells within a condition. Conditions are therefore paired within each repetition.

Table 9: CDC 2024 four-cell replication with the frozen 2022 pipelines. Capture changes and allocations are percentage points; capacity shares include central 95% ranges from 500 paired resamples.

Outcome	q	Joint gain	Eligibility	Capacity	Capacity share (%)
Preterm birth	5%	3.45	1.16	2.29	66.4 [64.9, 68.0]
Preterm birth	10%	5.85	2.23	3.62	61.9 [60.8, 63.0]
NICU admission	5%	3.34	1.15	2.19	65.6 [64.3, 67.0]
NICU admission	10%	5.66	2.11	3.55	62.8 [61.8, 63.9]
Low birth weight	5%	4.13	1.52	2.61	63.1 [61.7, 64.6]
Low birth weight	10%	7.01	2.74	4.27	60.9 [59.8, 61.9]

Table 10: Availability intercepts used in every synthetic condition. Both nonzero slopes have the same intercept by normal symmetry.

Target availability a	$\gamma = 0$	$\gamma = \pm 1$
0.5	0.0000000000	0.0000000000
0.7	0.8472978604	1.0184006520
0.9	2.1972245773	2.5642215001

Selection and reported uncertainty. Budgets use Python’s round of q times the relevant pool size. Descending score order breaks exact ties by ascending original row position. For the transported-threshold policy, the independent threshold sample uses $R^{\text{cal}} = Z^{\text{cal}} + \sigma\epsilon^{\text{cal}}$; its $\text{round}(q/50,000)$ -th highest score is applied with $\geq$. Every policy uses the same realized full-population event denominator within a repetition. Condition means and 2.5th/97.5th percentiles summarize 100 repetitions. These percentiles describe variation across generated populations, rather than a confidence interval for the Monte Carlo mean. The 18 sign reversals compare condition means; the nine negative ranges refer specifically to the all-record-budget eligibility contrast $C_{11} - C_{01}$. The 81 grid settings carry equal weight in the reported 22% reversal fraction, which describes this stress-test design rather than a prevalence of failure in clinical deployments.

The score-noise levels $\sigma = 0.5, 1, 2$ imply expected correlations between score and latent risk of approximately 0.894, 0.707, and 0.447. The largest absolute replicate-level interaction is 10.72 percentage points. Full condition summaries and replicate outputs are retained in the prepared release bundle.

D Audit Identities and Artifacts

Deterministic identity checks. For fixed S , scores, and deterministic tie breaking, $\text{Top}(S, b_e) \subseteq \text{Top}(S, b_a)$ whenever $b_e \leq b_a$. Increasing the budget adds exactly $b_a - b_e$ ranked positions and between zero and $b_a - b_e$ events. Both marginal capacity contrasts are therefore non-negative, as is their average ϕ_B . Substituting Equation 6 also gives $\phi_E + \phi_B = C_{11} - C_{00}$ exactly. These statements concern a shared score vector; retraining introduces an additional factor.

Artifacts. All manuscript figures and tables are generated from included CSV files. Verification scripts assert the six primary Shapley rows, 500 paired empirical resamples per target and budget, exact accounting identities, all 81 synthetic conditions, 18 mean sign reversals, and 9 conditions whose central replicate range is negative. The prepared release bundle contains aggregate inputs, figure and table generators, verification scripts, the locked synthetic configuration and runner, synthetic replicate outputs, and the corrected CDC 2024 replication aggregates. Local fitting and temporal-audit entrypoints are provided in the repository. The package redistributes no individual-level Natality records. Use of the public-use files follows the NCHS Data User Agreement [National Center for Health Statistics, 2025]; LightGBM is distributed under the MIT License [Microsoft Corporation and LightGBM Contributors, 2026].

Raw-data reconstruction. The CDC 2022 split is drawn before target-specific exclusions with `default_rng(2026)`: 20% enters validation and the remainder enters training, on which preprocessing is fitted. Unknown outcomes are excluded target-wise, and original row order breaks score ties. Thresholds are selected on the target-specific CDC 2022 all-record validation pool and applied with $\geq$. Empirical and synthetic seeds are 2027 and 2028. Section 3 and Appendices A–C specify the remaining cohort, model, and estimand details.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Sections 1 and 4 state the decision-audit contribution, empirical magnitude, and synthetic sign-reversal result using the same estimands as the Methods.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Sections 3, 5, and 6 and Appendix B state the population, score, operating-point, care-entry proxy, and conditional-resampling scope and identify the next validation tier.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper makes no standalone theoretical contribution. Section 3 defines the evaluation estimands, and Appendix D records elementary identities for frozen rankings and nested candidate sets.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 and Appendices A–D disclose the public data releases, cohort and outcome rules, split algorithm and seeds, features, preprocessing, model parameters, operating policies, tie breaking, resampling design, and full synthetic grid. The prepared artifact bundle regenerates every figure and table and verifies the accounting identities from archived aggregate inputs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: CDC Natality public-use files and official layouts are openly accessible. Code, aggregate artifacts, configurations, and reproduction instructions will be publicly released with the camera-ready version at <https://github.com/vache4ogan/prenatal-risk-scoreability>. The source package already contains the manuscript generators, synthetic experiment, and verification scripts. Repository access is pending the owner's visibility change.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 and Appendices A–D specify cohorts, outcomes, features, preprocessing, split algorithms and seeds, model settings, operating points, uncertainty, and every synthetic factor.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 3.4 defines 500 paired full-size record-level resamples with frozen scores and observed budgets. Appendix B reports their percentile ranges. Appendix C defines the synthetic central 95% replicate ranges, which measure variation across generated populations rather than uncertainty in their Monte Carlo mean. Figure 3 reports reversal frequencies. Neither experiment’s ranges include model-retraining uncertainty.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All runs are CPU-only; a multicore worker with 16 GB RAM is sufficient. Archived wall times are approximately 26 minutes for the canonical analysis, 2.5 minutes for the four-cell bootstrap, 18–21 minutes per feature ablation, and 20 seconds for the synthetic grid.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The study uses public-use deidentified vital-statistics records, releases no individual records or deployable clinical model, and frames results as evaluation evidence.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 5 discusses transparent resource accounting as a benefit and hidden workload expansion and redirected scarce services as operational risks.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work releases neither a deployable clinical model nor a new sensitive dataset.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Section 3 and Appendix D cite the NCHS Data User Agreement and the LightGBM MIT License. The package redistributes only derived aggregate results.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The prepared release bundle documents the audit inputs and outputs, exact commands, environment, aggregate source-data schema, and controlled-experiment code. It contains no individual-level dataset or pretrained clinical model.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study includes no crowdsourcing or direct participant interaction.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The analysis uses deidentified public-use records and introduces no participant recruitment or intervention.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an experimental or methodological component of the study.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.