
Same Top Fraction, Different Workload: Auditing Eligibility and Capacity in Clinical Risk Evaluation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Clinical risk scores turn patient rankings into finite worklists for specialist review,
2 monitoring, or prevention. Yet cross-cohort evaluations often compare the same
3 top fraction across candidate sets of different sizes, changing both who can be
4 ranked and how many people receive service. We introduce a decision-aware four-
5 cell audit that keeps the full-population event denominator fixed and crosses two
6 eligibility rules with two absolute budgets, separating broader reach from addi-
7 tional capacity. We apply it to frozen birth-record score vectors for preterm birth,
8 neonatal intensive care admission, and low birth weight in approximately 3.48 mil-
9 lion CDC Natality 2023 records. At the nominal top 10% operating point, moving
10 from early-entry eligibility to a retrospective all-record upper bound reports a 39–
11 42% relative increase in captured events while expanding the worklist by one third.
12 Symmetric accounting allocates about three fifths of the joint gain to added capac-
13 ity and two fifths to broader eligibility; at a fixed budget, the eligibility gain is
14 less than half the cohort-relative estimate. In 22% of controlled conditions, the
15 cohort-relative comparison improves while fixed-budget capture declines. The au-
16 dit therefore distinguishes population reach from resource expansion and makes
17 the policy represented by a top-fraction metric explicit.

18 1 Introduction

19 Clinical risk models increasingly determine which records receive specialist review, monitoring,
20 outreach, or preventive services [Annarumma et al., 2019, Baltruschat et al., 2021, Jung et al., 2021,
21 Singh et al., 2023]. These workflows are capacity constrained: a model produces a ranking, but an or-
22 ganization acts on a finite number of cases. Top- q precision, recall, or event capture therefore appears
23 to connect predictive performance directly to an operational decision. Prenatal case-management
24 studies, for example, evaluate the highest-risk 5–15% because program capacity covers a fraction of
25 eligible pregnancies [Pan et al., 2017].

26 The concrete scale of this policy difference is visible in CDC Natality 2023 [National Center for
27 Health Statistics, 2026]. An early-entry candidate set of approximately 2.60 million records yields a
28 top-10% worklist of about 260,000, whereas an all-record set of approximately 3.48 million yields
29 about 348,000 selections at the same nominal fraction. The comparison therefore changes both
30 the candidate gate and the absolute workload. Clinical decision frameworks distinguish predictive
31 ranking from action under finite resources [Jung et al., 2021, Singh et al., 2023]: a fixed-capacity
32 service and a service that expands with candidate volume implement different policies and require
33 different contrasts.




Figure 1: The same nominal top 10% encodes a different workload. Panel (a) summarizes the target-specific CDC 2023 cohorts: moving from the early-entry gate to the all-record upper bound increases the selected count from about 260,000 to 348,000. Panel (b) allocates the associated increase in population event capture at $q = 10\%$. Values inside bars are percentage-point allocations; capacity contributes about 62% of each total gain.

Prenatal risk triage makes this ambiguity visible. Pregnancies enter care at different times, and adverse outcomes are concentrated among records with no documented prenatal care or unknown start timing. An early-entry score ranks events inside its candidate gate; events outside the gate contribute zero reachable cases regardless of discrimination. National vital statistics retain both the eligible and excluded outcomes, making it possible to distinguish retrieval within the gate from reach across the complete target population.

We audit this distinction on CDC Natality preterm birth, neonatal intensive care unit (NICU) admission, and low birth weight. Figure 1 previews the result. At the nominal top 10% operating point, all-record upper-bound evaluation selects roughly 88,000 additional records per outcome. Population event capture rises substantially, but about 62% of the joint increase is allocated to the larger workload. A controlled mechanism grid then identifies conditions in which the resource change is large enough to reverse the fixed-budget conclusion.

Our contributions are:

1. **An auditable decision estimand.** Population event capture keeps every target-population event in the denominator, while a complete eligibility-by-budget design exposes the two fixed-budget contrasts and their interaction.
2. **A national-scale measurement study.** Frozen 2022 scoring pipelines applied to CDC 2023 show that the same top 10% expands workload by 33.8% and that capacity receives 61.6–62.2% of the symmetric allocation across three outcomes.
3. **A mechanism map and reporting profile.** An 81-condition stress test locates sign-reversal regimes, and the resulting checklist makes the candidate gate, absolute workload, and event denominator explicit.

Code for the four-cell audit and synthetic experiments, together with the aggregate artifacts used to generate the reported figures and tables, will be released upon publication.

2 Related Work

Risk ranking under limited resources. Prenatal case-management, imaging worklists, and clinical workflow studies connect risk scores to finite review or treatment capacity [Pan et al., 2017, Annarumma et al., 2019, Baltruschat et al., 2021, Jung et al., 2021, Singh et al., 2023]. They motivate operating points such as top 5%, top 10%, or a fixed number of patients and show why use-aligned metrics matter. Calibrated subset selection similarly constructs shortlists with guarantees inside a supplied candidate pool [Wang et al., 2022]. These approaches make selection within a declared pool more reliable, while leaving cross-pool comparisons dependent on how the shortlist size changes with the pool.

Selective access to clinical data and services. Clinical records reflect who enters care, is tested, and receives a documented label. Informed presence and selection bias can alter both apparent risk and the population represented by an EHR cohort [Haneuse and Daniels, 2016, Goldstein et al., 2016, Sisk et al., 2021]; disparate censorship describes related variation in testing and label availability [Chang et al., 2022]. Health-service coverage frameworks separate service availability from realized coverage [Tanahashi, 1978]. This literature establishes that the candidate gate is informative, but a top-fraction comparison adds a second moving quantity because the number served scales with the observed cohort.

Evaluation validity and decomposition. Clinical reporting guidance requires explicit target populations, intended use, discrimination, and calibration [Collins et al., 2024]. Symmetric allocation across factor orderings is established decomposition machinery [Kitagawa, 1955, Shapley, 1953, Shorrocks, 2013]. Taken together, these conventions still leave the operational meaning of a metric implicit when candidate membership and absolute capacity change together. The case is common and concrete: a percentage-based budget is held numerically constant across cohorts, its absolute size moves, and the resulting gain is reported as if it represented one intervention.

3 Methodology

The audit fixes the target population, score vector, and event denominator, then crosses two candidate-set rules with two absolute budgets. Figure 2 shows the resulting four cells and the two paths whose average yields the symmetric allocation.




Figure 2: Eligibility–capacity audit. Rows vary the candidate set and columns vary the absolute selection budget while the score vector and full-population event denominator remain fixed. Blue arrows change candidate eligibility; orange arrows change capacity. Averaging the two paths from C_{00} to C_{11} gives the symmetric Shapley allocation, whereas $C_{11} - C_{00}$ is the usual same- q comparison and changes both factors.

3.1 Target population and scoring study

We use CDC Natality public-use live-birth records and official field definitions [National Center for Health Statistics, 2026, 2017]. CDC 2022 supplies model development with an 80/20 train/validation split; the frozen pipelines generate the shared score vectors analyzed in CDC 2023. Model fitting and hyperparameter selection use CDC 2022 exclusively; CDC 2023 is the temporal cohort in which the audit hypothesis was developed and evaluated. The target population contains U.S.-resident singleton live births with a known outcome: preterm birth below 37 completed obstetric-estimate weeks ($N = 3,480,382$, 303,001 events), NICU admission ($N = 3,476,584$, 303,285 events), and birth weight below 2,500 g ($N = 3,480,316$, 243,776 events).

Prenatal-care month defines four mutually exclusive groups: early entry in months 1–3, later documented entry in months 4–10, no prenatal care, and unknown start timing. The early candidate set contains months 1–3. The documented-start set contains months 1–10. The all-record set adds no-care and unknown-timing records and therefore bounds candidate and event availability under retrospective complete scoring. The care-month rule operationalizes documented entry into prenatal

care rather than feature-level EHR availability; the birth-certificate model serves as a shared retrospective ranking instrument.

The baseline-oriented score uses nine maternal birth-certificate variables: race, height, pre-pregnancy BMI and weight, previous live births now living, previous live births now dead, other pregnancy outcomes, pre-pregnancy diabetes, and pre-pregnancy hypertension. Outcome, delivery, and neonatal fields are excluded. Numeric variables use development-set median imputation and standardization; race uses most-frequent imputation and one-hot encoding. One LightGBM model per outcome uses 350 trees, learning rate 0.05, 31 leaves, minimum child size 100, and seed 2026 [Ke et al., 2017]. The audit compares candidate and budget policies on a shared score vector while holding the model family constant.

3.2 Decision metrics

Let $\mathcal{P}$ contain N records and $E = \sum_{i \in \mathcal{P}} Y_i$ events. For candidate set S , event count E_S , and integer budget b , deterministic tie breaking defines the b highest scores $\text{Top}(S, b)$. Their true-positive count is

$$\text{TP}(S, b) = \sum_{i \in \text{Top}(S, b)} Y_i. \quad (1)$$

We report both within-gate performance and population reach:

$$\text{Precision}(S, b) = \frac{\text{TP}(S, b)}{b}, \quad \text{Recall}_{\text{eligible}}(S, b) = \frac{\text{TP}(S, b)}{E_S}, \quad (2)$$

$$\text{Capture}(S, b) = \frac{\text{TP}(S, b)}{E}, \quad \text{Ceiling}(S) = \frac{E_S}{E}. \quad (3)$$

Population capture factors into retrieval and reach,

$$\text{Capture}(S, b) = \text{Recall}_{\text{eligible}}(S, b) \text{ Ceiling}(S). \quad (4)$$

Each result therefore reports target-population N and E , eligible n and E_S , exact selected n , precision, eligible recall, population capture, and the event ceiling. ROC-AUC and average precision (AP) provide a ranking-signal check; the four-cell contrasts carry the primary decision claim.

3.3 Four-cell eligibility–capacity audit

Let $S_e \subseteq S_a$ be early and all-record candidate sets. A nominal fraction q induces $b_e = \text{round}(q|S_e|)$ and $b_a = \text{round}(q|S_a|)$. With $b_a \leq |S_e|$, the complete design is

$$C_{uv} = \text{Capture}(S_u, b_v), \quad u, v \in \{0, 1\}, \quad (5)$$

where 0 denotes early and 1 denotes all-record eligibility or capacity. The usual cohort-relative contrast is $\Delta = C_{11} - C_{00}$. Its interaction is $I = C_{11} - C_{10} - C_{01} + C_{00}$. We retain both fixed-budget eligibility contrasts and summarize the joint change with the standard symmetric Shapley allocation [Shapley, 1953]:

$$\begin{aligned} \phi_E &= \frac{1}{2} [(C_{10} - C_{00}) + (C_{11} - C_{01})], \\ \phi_B &= \frac{1}{2} [(C_{01} - C_{00}) + (C_{11} - C_{10})]. \end{aligned} \quad (6)$$

The exact identity $\Delta = \phi_E + \phi_B$ holds, and $s_B = \phi_B/\Delta$ gives the symmetric capacity share for positive Δ . For a frozen deterministic ranking, increasing b adds nested positions, so the capacity component is non-negative; expanding S at fixed b can have either sign. A uniformly random ranking supplies the mechanical reference $E[\text{Capture}(S, b)] = (b/|S|)(E_S/E)$.

3.4 Evaluation and mechanism stress test

The empirical audit uses $q \in \{0.05, 0.10\}$. Point estimates are accompanied by 500 paired full-size empirical resamples per outcome and q . Each resample uses common record multiplicities across all four cells while holding the observed scores and absolute budgets fixed. The resulting central 95% ranges measure conditional stability of the reported four-cell contrasts.

The controlled simulation crosses availability fractions 0.5, 0.7, and 0.9; risk–availability relationships in which higher risk is less available, independent, or more available; score noise levels 0.5, 1,

and 2; and $q \in \{0.05, 0.10, 0.20\}$. These factors yield 81 conditions. Each of 100 repetitions contains 100,000 evaluation records and an independent 50,000-record threshold-selection population. Outcomes, availability, and scores are shared across the compared policies within a repetition.

4 Results

4.1 Candidate eligibility and event ceilings

Early entry contains about 74.7% of target-population records, but only 71.8% of preterm events, 70.6% of NICU events, and 69.4% of low-birth-weight events. Before ranking begins, the early gate therefore leaves 28.2–30.6% of events outside its reachable population. Including every documented prenatal-care start in months 1–10 raises the event ceiling to 91.1–92.1%.

In absolute terms, early eligibility excludes 74,624–89,111 outcome events across the three tasks. Even the documented-start gate leaves 20,558–27,087 events outside the candidate population. These are reach constraints: changing the ranking leaves events outside its decision set unreachable.

The excluded records carry concentrated risk. Outcome prevalence among no-care records is 2.4–2.7 times the early-entry prevalence across the three targets; unknown-timing records are 1.7–1.9 times higher. Candidate expansion therefore changes both population reach and case mix, making a full-population event denominator essential.

Table 1: The candidate gate fixes reach before ranking. Early entry denotes prenatal-care initiation in months 1–3; documented start includes months 1–10. Risk ratios compare outcome prevalence in excluded groups with the early-entry group.

Outcome	Early eligible n	Early share	Early ceiling	Documented ceiling	No-care / unknown risk
Preterm birth	2,601,060	74.7%	71.8%	91.1%	$2.7 \times / 1.8 \times$
NICU admission	2,597,274	74.7%	70.6%	92.1%	$2.4 \times / 1.7 \times$
Low birth weight	2,600,093	74.7%	69.4%	91.6%	$2.6 \times / 1.8 \times$

4.2 Eligibility–capacity accounting under top- q selection

Figure 1 and Table 2 connect the nominal operating point to its actual workload. Across outcomes, the selected count increases from approximately 260,000 to 348,000. Population event capture rises from 14.63% to 20.30% for preterm birth, 13.93% to 19.52% for NICU, and 16.06% to 22.83% for low birth weight. These are relative gains of 38.8%, 40.1%, and 42.2%, respectively.

The proportional policy adds 87,931–88,023 service slots per outcome and captures 16,505–17,177 additional events. Its marginal yield is therefore 18.8–19.5 additional events per 100 added slots, equivalent to 5.1–5.3 additional selections per extra captured event. Thus, the higher capture is obtained under an explicitly 33.8% larger service capacity.

Table 2: Same nominal top 10%, different operational result. Population capture uses the full target-population event denominator. ϕ_E and ϕ_B are the symmetric eligibility and capacity allocations; conditional empirical-resampling ranges appear in Appendix 7.

Outcome	Selected n	Workload	Capture (%)	Rel. gain	ϕ_E / ϕ_B (pp)	s_B
Preterm birth	260,106→348,038	+33.8%	14.63→20.30%	+38.8%	2.15/ 3.52	62.1%
NICU admission	259,727→347,658	+33.9%	13.93→19.52%	+40.1%	2.11/ 3.47	62.2%
Low birth weight	260,009→348,032	+33.9%	16.06→22.83%	+42.2%	2.60/ 4.17	61.6%

The four-cell audit changes the interpretation of those gains. Symmetric accounting assigns 2.11–2.60 percentage points to eligibility and 3.47–4.17 points to additional capacity, placing 61.6–62.2% of the joint contrast on the workload increase. Holding the early budget fixed yields eligibility gains of 1.89–2.27 points; holding the all-record budget fixed yields 2.32–2.93 points. Both are materially smaller than the joint 5.59–6.77-point cohort-relative contrast.

The pattern persists at $q = 5\%$, where capacity receives 63.9–65.6% of the symmetric allocation. Conditional ranges remain above 60.4% for every target and operating point. A random ranking allocates 79.7–87.9% to capacity at $q = 10\%$, showing that a majority capacity allocation can arise mechanically under a proportional budget. The learned ranking shifts more of the joint gain toward eligibility because the added candidate groups contain greater event mass. Complete four-cell counts and both operating points appear in Appendix B.

4.3 Fixed-budget and transported-threshold policies

Table 3 compares three policies on the same frozen score vector. Cohort-relative top-10% evaluation jointly changes eligibility and workload. Fixing the all-record absolute budget isolates candidate expansion: capture increases by 2.32–2.93 points, with exactly the same number selected on both sides. A threshold chosen once on the CDC 2022 validation split selects 9.15–9.82% of early-entry records and 10.20–10.40% of all records in CDC 2023. Under that transported threshold, the all-record workload is 1.42–1.49 times larger and capture increases by 6.36–8.19 points.

Table 3: One score vector, three resource policies at the nominal 10% operating point. Counts and population capture compare early-entry with all-record eligibility. The score threshold is chosen on CDC 2022 and transported unchanged to CDC 2023.

Outcome	Resource policy	Selected n : early $\rightarrow$ all	Capture: early $\rightarrow$ all	Δ (pp)
Preterm birth	Cohort-relative top 10%	260,106 $\rightarrow$ 348,038	14.63 $\rightarrow$ 20.30%	+5.67
	Fixed absolute b_a	348,038 $\rightarrow$ 348,038	17.89 $\rightarrow$ 20.30%	+2.40
	Transported 2022 threshold	245,487 $\rightarrow$ 361,709	14.06 $\rightarrow$ 20.84%	+6.78
NICU admission	Cohort-relative top 10%	259,727 $\rightarrow$ 347,658	13.93 $\rightarrow$ 19.52%	+5.59
	Fixed absolute b_a	347,658 $\rightarrow$ 347,658	17.20 $\rightarrow$ 19.52%	+2.32
	Transported 2022 threshold	255,166 $\rightarrow$ 361,714	13.76 $\rightarrow$ 20.12%	+6.36
Low birth weight	Cohort-relative top 10%	260,009 $\rightarrow$ 348,032	16.06 $\rightarrow$ 22.83%	+6.77
	Fixed absolute b_a	348,032 $\rightarrow$ 348,032	19.90 $\rightarrow$ 22.83%	+2.93
	Transported 2022 threshold	237,973 $\rightarrow$ 355,034	14.98 $\rightarrow$ 23.17%	+8.19

These contrasts differ because they implement different resource commitments on a shared ranking. A complete cross-cohort claim therefore states its candidate rule, selected count, and event denominator alongside the gain.

The cross-year workload check points in the same direction. Transporting the stored CDC 2022 top-10% score cutoffs to target-specific CDC 2024 cohorts selected 10.65% of preterm records, 9.50% of NICU records, and 8.30% of low-birth-weight records. Relative to a 10% service budget, these rates create 22,668 additional preterm slots but leave 17,621 NICU and 59,716 low-birth-weight slots unfilled. A fixed score cutoff therefore defines a workload that can move across years; the exact selected count remains the operational quantity for capacity planning.

4.4 Controlled mechanism stress test

Figure 3 maps the protocol difference across all 81 synthetic conditions. Eighteen conditions report a positive cohort-relative gain while candidate expansion at the all-record absolute budget reduces population capture. Nine have a central replicate range entirely below zero. Every reversal arises when higher-risk records are preferentially available, so the newly admitted pool is lower risk and can displace selected events.

At 70% availability, medium score quality, and $q = 10\%$, for example, cohort-relative capture increases 5.52 points while fixed-budget candidate expansion changes capture by -0.28 points; 93 of 100 repetitions reverse. The largest mean absolute interaction is 9.49 points. In CDC 2023, all six fixed-budget eligibility contrasts remain positive, so the ambiguity changes effect magnitude. The controlled grid shows that selective availability can also change effect direction.

4.5 Temporal score diagnostics

The four-cell accounting fixes one CDC 2023 score vector, isolating its within-year policy contrasts from score drift. A temporal check additionally shows that the frozen 2022 pipelines retain similar




Figure 3: Full-grid mechanism stress test. Panel (a) compares the cohort-relative conclusion with candidate expansion at the all-record-derived absolute budget b_a . The shaded quadrant contains sign reversals: the joint top- q contrast is positive while fixed-budget eligibility expansion is negative. Panel (b) gives replicate-level reversal frequency for conditions where higher risk is more available. Marker shape denotes score quality.

aggregate discrimination in CDC 2023 and provide a transported ranking for the policy audit. Across outcomes, ROC-AUC changes by at most 0.003 and average precision by at most 0.003 from the CDC 2022 validation split to CDC 2023. CDC 2023 Brier skill relative to a prevalence-only forecast is 1.7–2.4%. These diagnostics anchor the empirical audit in transported scores while preserving the decision policy as the primary estimand.

Table 4: Frozen-score temporal check. Brier baseline is the constant CDC 2023 prevalence prediction; skill is $1 - \text{Brier}_{\text{model}} / \text{Brier}_{\text{baseline}}$.

Outcome	ROC-AUC: 2022 $\rightarrow$ 2023	AP: 2022 $\rightarrow$ 2023	Brier: model / baseline	Brier skill
Preterm birth	0.614 $\rightarrow$ 0.612	0.146 $\rightarrow$ 0.146	0.0780 / 0.0795	1.9%
NICU admission	0.611 $\rightarrow$ 0.608	0.139 $\rightarrow$ 0.141	0.0783 / 0.0796	1.7%
Low birth weight	0.653 $\rightarrow$ 0.652	0.133 $\rightarrow$ 0.132	0.0636 / 0.0651	2.4%

4.6 Sensitivity and robustness

Two feature ablations preserve the allocation pattern. Removing maternal race and refitting each outcome model at $q = 10\%$ leaves capacity shares of 60.9–62.4%. A separate eight-feature specification removes prior other pregnancy outcomes (prior_terminations) and repeats the complete audit at both budgets with paired resampling. Capacity still accounts for 61.5–64.3% of the joint contrast, and all interaction ranges remain above zero (+0.32 to +0.69 percentage points). The result therefore persists across two alternative conservative landmark specifications.

5 Discussion

Resource interpretation of top- q evaluation. Top- q evaluation is operationally coherent when service volume scales with the candidate population; a fixed absolute budget answers the constrained capacity question. The distinction is material here: the headline 38.8–42.2% relative capture gain includes a 33.8% workload increase. The four-cell design makes that resource policy visible and supplies both fixed-budget contrasts.

Separating reach, ranking, and capacity. The candidate gate determines the event ceiling, ranking determines which eligible cases approach the cutoff, and capacity determines how many positions can be filled. Freezing the score vector isolates broader reach at fixed capacity from proportional resource expansion; the interaction records when the value of one lever depends on the other.

Transported score cutoffs form a third regime in which workload emerges from the subsequent score distribution, as the CDC 2024 check illustrates.

Applicability across deterministic score vectors. The four-cell accounting accepts any fixed deterministic score vector, and capacity monotonicity follows from the nesting of top- b selections independently of the LightGBM functional form. Candidate expansion at fixed b depends on which scores and events enter the pool, so its sign remains empirical. The random-ranking reference and the controlled grid spanning three score-noise levels recover both properties: capacity usually explains most of a proportional-budget contrast, while selective availability can reverse the fixed-budget effect. Model choice changes where a study lies in this mechanism map; the audit defines the decision comparison around it.

A recommended reporting profile. Table 5 turns the audit into a compact submission and deployment reporting template. These fields complement clinical prediction guidance by connecting a cohort-level score to the exact decision implemented at its operating point.

Table 5: Recommended reporting profile for cross-cohort top- q evaluation. Each item identifies one part of the operational claim.

Dimension	Report	Question answered
Target population	Population rule, N , total events E	Whose outcomes define success?
Candidate gate	Eligibility rule, eligible n , E_S , event ceiling	Who can be ranked and reached?
Resource rule	Fraction q or budget b , with exact selected n	How much service is available?
Operating result	Precision, eligible recall, population capture	What is retrieved inside and beyond the gate?
Cross-cohort audit	Four cells, fixed-budget contrasts, interaction	Which policy change produced the gain?

The profile applies wherever a score becomes a finite worklist: imaging queues, screening, case management, and public-sector allocation all require a declared fixed, proportional, or threshold-driven capacity rule. Exact selected counts prevent a stable percentage from disguising a changing workload.

Implementation in existing evaluation pipelines. Stored identifiers, outcomes, candidate masks, and one score vector are enough to construct the four cells. Implementations record both budgets, selected events, and the full-population denominator, then reuse common record weights for paired resampling. Aggregate counts and configurations provide an auditable path from a top- q result to its service policy without releasing row-level clinical data.

6 Conclusion

Top-fraction comparisons across cohorts encode service decisions in addition to ranking summaries. In the CDC audit, the same nominal top 10% expands the workload by one third, and symmetric accounting assigns about 62% of the capture gain to that capacity increase. Controlled experiments identify conditions in which cohort-relative and fixed-budget conclusions have opposite signs. Reporting the target population, candidate rule, exact selected count, full-population event denominator, and fixed-budget contrast makes the policy claim auditable. The present scope is a birth-certificate score and a care-entry eligibility proxy; prospective EHR landmarks, subgroup budgets, and additional temporal cohorts provide the next transport tests.

References

- Mauro Annarumma, Samuel J. Withey, Robert J. Bakewell, Emanuele Pesce, Vicky Goh, and Giovanni Montana. Automated triaging of adult chest radiographs with deep artificial neural networks. *Radiology*, 291(1):196–202, 2019. doi: 10.1148/radiol.2018180921.
- Ivo M. Baltruschat, Leonhard Steinmeister, Hannes Nickisch, Axel Saalbach, Michael Grass, Gerhard Adam, Tobias Knopp, and Harald Ittrich. Smart chest X-ray worklist prioritization using artificial intelligence: A clinical workflow simulation. *European Radiology*, 31(6):3837–3845, 2021. doi: 10.1007/s00330-020-07480-7.
- Trenton Chang, Michael W. Sjoding, and Jenna Wiens. Disparate censorship and undertesting: A source of label bias in clinical machine learning. In *Proceedings of the 7th Machine Learning for Healthcare Conference*, volume 182 of *Proceedings of Machine Learning Research*, pages 343–390, 2022. URL <https://proceedings.mlr.press/v182/chang22a.html>.
- Gary S. Collins, Karel G. M. Moons, Paula Dhiman, Richard D. Riley, Andrew L. Beam, Ben Van Calster, Marzyeh Ghassemi, Xiaoxuan Liu, Johannes B. Reitsma, Maarten van Smeden, et al. TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods. *BMJ*, 385:e078378, 2024. doi: 10.1136/bmj-2023-078378.
- Benjamin A. Goldstein, Nrupen A. Bhavsar, Matthew Phelan, and Michael J. Pencina. Controlling for informed presence bias due to the number of health encounters in an electronic health record. *American Journal of Epidemiology*, 184(11):847–855, 2016. doi: 10.1093/aje/kww112.
- Sebastien Haneuse and Michael J. Daniels. A general framework for considering selection bias in EHR-based studies: What data are observed and why? *eGEMs*, 4(1):1203, 2016. doi: 10.13063/2327-9214.1203.
- Kenneth Jung, Sehj Kashyap, Anand Avati, Stephanie Harman, Heather Shaw, Ron C. Li, Margaret Smith, Kenny Shum, Jacob Javitz, Yohan Vetteth, Tina Seto, Steven C. Bagley, and Nigam H. Shah. A framework for making predictive models useful in practice. *Journal of the American Medical Informatics Association*, 28(6):1149–1158, 2021. doi: 10.1093/jamia/ocaa318.
- Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL https://proceedings.neurips.cc/paper_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html.
- Evelyn M. Kitagawa. Components of a difference between two rates. *Journal of the American Statistical Association*, 50(272):1168–1194, 1955. doi: 10.1080/01621459.1955.10501299.
- Microsoft Corporation and LightGBM Contributors. Lightgbm. GitHub repository, MIT License, 2026. URL <https://github.com/lightgbm-org/LightGBM>. Accessed 2026-08-26.
- National Center for Health Statistics. Facility worksheet for the live birth certificate. Centers for Disease Control and Prevention, 2017. URL <https://www.cdc.gov/nchs/data/dvs/facility-worksheet-2016.pdf>. Final 12/16 version; accessed 2026-08-26.
- National Center for Health Statistics. Nchs data user agreement. Centers for Disease Control and Prevention, 2025. URL <https://www.cdc.gov/nchs/policy/data-user-agreement.html>. Accessed 2026-08-26.
- National Center for Health Statistics. Vital statistics online data portal: Natality public-use data files and documentation. Centers for Disease Control and Prevention, 2026. URL https://www.cdc.gov/nchs/data_access/vitalstatsonline.htm. Accessed 2026-08-26.
- Ian Pan, Laura B. Nolan, Rashida R. Brown, Romana Khan, Paul van der Boor, Daniel G. Harris, and Rayid Ghani. Machine learning for social services: A study of prenatal case management in Illinois. *American Journal of Public Health*, 107(6):938–944, 2017. doi: 10.2105/AJPH.2017.303711.

- 305 Lloyd S. Shapley. A value for n -person games. In Harold W. Kuhn and Albert W. Tucker, editors,
 306 *Contributions to the Theory of Games II*, pages 307–317. Princeton University Press, Princeton,
 307 1953. doi: 10.1515/9781400829156-012.
- 308 Anthony F. Shorrocks. Decomposition procedures for distributional analysis: A unified framework
 309 based on the shapley value. *The Journal of Economic Inequality*, 11(1):99–126, 2013. doi: 10.1
 310 007/s10888-011-9214-z.
- 311 Karandeep Singh, Nigam H. Shah, and Andrew J. Vickers. Assessing the net benefit of machine
 312 learning models in the presence of resource constraints. *Journal of the American Medical Informatics Association*, 30(4):668–673, 2023. doi: 10.1093/jamia/ocad006.
- 314 Rose Sisk, Lijing Lin, Matthew Sperrin, Jessica K. Barrett, Brian Tom, Karla Diaz-Ordaz, Niels
 315 Peek, and Glen P. Martin. Informative presence and observation in routine health data: A re-
 316 view of methodology for clinical risk prediction. *Journal of the American Medical Informatics Association*, 28(1):155–166, 2021. doi: 10.1093/jamia/ocaa242.
- 318 T. Tanahashi. Health service coverage and its evaluation. *Bulletin of the World Health Organization*,
 319 56(2):295–303, 1978. URL <https://pmc.ncbi.nlm.nih.gov/articles/PMC2395571/>.
- 320 Lequn Wang, Thorsten Joachims, and Manuel Gomez Rodriguez. Improving screening processes
 321 via calibrated subset selection. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 22702–22726, 2022.
 322 URL <https://proceedings.mlr.press/v162/wang22j.html>.
 323

324 A Outcome, Model, and Cohort Details

Table 6: Baseline-oriented birth-certificate variables used by the frozen ranking instrument. Encounter-specific availability at months 1–3 requires validation in a prospective EHR source.

Feature	CDC source meaning
Maternal race	Maternal characteristic on birth record
Maternal height	Maternal characteristic on birth record
Pre-pregnancy BMI	Derived from pre-pregnancy measures
Pre-pregnancy weight	Weight before pregnancy
Previous live births now living	Previous live-born infants still living
Previous live births now dead	Previous live-born infants no longer living
Other pregnancy outcomes	Non-live-birth outcomes; may include current pregnancy
Pre-pregnancy diabetes	Diabetes diagnosed before this pregnancy
Pre-pregnancy hypertension	Hypertension diagnosed before this pregnancy

325 For CDC 2023 preterm records, early entry contains 2,601,060 records and 217,632 events; later
 326 entry contains 739,298 and 58,282; no care contains 76,520 and 17,387; unknown timing contains
 327 63,504 and 9,700. Target-specific counts differ slightly because each outcome has its own missing-
 328 label exclusions. Source code 99 is normalized to missing before cohort construction, while source
 329 value 0 denotes no prenatal care.

330 B Complete Four-Cell Results

331 The point-estimate budgets at $q = 5\%$ are 130,053/174,019 for preterm, 129,864/173,829 for NICU,
 332 and 130,005/174,016 for low birth weight (early/all-record). At $q = 10\%$ they are 260,106/348,038,
 333 259,727/347,658, and 260,009/348,032. These observed integer budgets remain fixed across empir-
 334 ical resamples.

335 C Synthetic Grid Details

336 The synthetic experiment crosses three availability fractions, three risk–availability slopes, three
 337 score-noise levels, and three selection fractions. Intercepts are solved numerically so the target

Table 7: Complete symmetric-allocation results. All effects are percentage-point changes in population event capture; brackets give central 95% paired empirical-resampling ranges conditional on frozen scores and observed budgets.

Outcome	q	Joint Δ	ϕ_E	ϕ_B	Symmetric share s_B	Interaction
Preterm birth	5%	3.37 [3.30, 3.43]	1.16 [1.09, 1.22]	2.21 [2.17, 2.25]	65.6 [64.3, 67.1]%	0.33 [0.26, 0.41]
Preterm birth	10%	5.67 [5.58, 5.77]	2.15 [2.06, 2.24]	3.52 [3.48, 3.57]	62.1 [61.1, 63.3]%	0.51 [0.44, 0.60]
NICU admission	5%	3.31 [3.26, 3.38]	1.19 [1.13, 1.25]	2.12 [2.09, 2.16]	64.1 [62.8, 65.4]%	0.29 [0.23, 0.36]
NICU admission	10%	5.59 [5.51, 5.67]	2.11 [2.03, 2.20]	3.47 [3.44, 3.52]	62.2 [61.1, 63.3]%	0.42 [0.34, 0.51]
Low birth weight	5%	3.96 [3.89, 4.06]	1.43 [1.34, 1.52]	2.53 [2.50, 2.57]	63.9 [62.3, 65.4]%	0.40 [0.32, 0.49]
Low birth weight	10%	6.77 [6.67, 6.89]	2.60 [2.49, 2.72]	4.17 [4.12, 4.22]	61.6 [60.4, 62.7]%	0.66 [0.55, 0.78]

Table 8: Auditable four-cell point estimates. Each cell reports population event capture in percent followed by selected true events in parentheses. The two budgets are exact selected counts.

Outcome	q	Budgets b_e/b_a	$C_{00}: S_e, b_e$	$C_{01}: S_e, b_a$	$C_{10}: S_a, b_e$	$C_{11}: S_a, b_a$
Preterm birth	5%	130,053/174,019	9.05 (27,434)	11.10 (33,618)	10.04 (30,432)	12.42 (37,631)
Preterm birth	10%	260,106/348,038	14.63 (44,326)	17.89 (54,216)	16.52 (50,061)	20.30 (61,503)
NICU admission	5%	129,864/173,829	8.51 (25,801)	10.48 (31,795)	9.55 (28,967)	11.82 (35,854)
NICU admission	10%	259,727/347,658	13.93 (42,252)	17.20 (52,160)	15.84 (48,033)	19.52 (59,201)
Low birth weight	5%	130,005/174,016	9.53 (23,229)	11.86 (28,918)	10.76 (26,224)	13.49 (32,892)
Low birth weight	10%	260,009/348,032	16.06 (39,145)	19.90 (48,505)	18.33 (44,679)	22.83 (55,650)

338 marginal prevalence or availability fraction holds before Monte Carlo variation. Common random
339 numbers are used across all 81 conditions within each repetition, and the available-subset and all-
340 population policies share outcomes and score vectors. Reported 2.5th and 97.5th percentiles describe
341 the replicate distribution.

342 The score-noise levels $\sigma_R = 0.5, 1, 2$ imply expected correlations between score and latent risk
343 of approximately 0.894, 0.707, and 0.447. The largest absolute replicate-level interaction is 10.72
344 percentage points. Full condition summaries and replicate outputs accompany the review package.

345 D Audit Identities and Artifacts

346 **Deterministic identity checks.** For fixed S , scores, and deterministic tie breaking, $\text{Top}(S, b_e) \subseteq$
347 $\text{Top}(S, b_a)$ whenever $b_e \leq b_a$. Increasing the budget adds exactly $b_a - b_e$ ranked positions and
348 between zero and $b_a - b_e$ events. Both marginal capacity contrasts are therefore non-negative, as
349 is their average ϕ_B . Substituting Equation 6 also gives $\phi_E + \phi_B = C_{11} - C_{00}$ exactly. These
350 statements concern a shared score vector; retraining introduces an additional factor.

351 **Artifacts.** All manuscript figures and tables are generated from included CSV files. Verification
352 scripts assert the six primary Shapley rows, 500 paired empirical resamples per target and budget,
353 exact accounting identities, all 81 synthetic conditions, 18 mean sign reversals, and 9 conditions
354 whose central replicate range is negative. The package includes aggregate inputs, generators, locked
355 analysis configurations, empirical runners, and synthetic replicate outputs. CDC reruns additionally
356 require the public-use Natality files and fitted model artifacts identified by the run manifests. The
357 package redistributes no individual-level Natality records. Use of the public-use files follows the
358 NCHS Data User Agreement [National Center for Health Statistics, 2025]; LightGBM is distributed
359 under the MIT License [Microsoft Corporation and LightGBM Contributors, 2026].

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Sections 1 and 4 state the decision-audit contribution, empirical magnitude, and synthetic sign-reversal result using the same estimands as the Methods.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 6 states the population, score, operating-point, care-entry proxy, and conditional-resampling scope and identifies the next validation tier.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper makes no standalone theoretical contribution. Section 3 defines the evaluation estimands, and Appendix D records elementary identities for frozen rankings and nested candidate sets.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: The anonymous package reproduces every manuscript figure and table, the accounting identities, and the synthetic study from archived aggregate artifacts. Reconstructing the frozen CDC scores from raw public-use files additionally requires the finalized extraction lineage and fitted-model bundle planned for the archival release.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility.

In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: CDC Natalty is publicly accessible and the blind-review package contains analysis code and aggregate artifacts. No anonymous public repository is provided; code and derived artifacts will be released upon publication.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 and Appendices A–C specify cohorts, outcomes, features, preprocessing, model settings, operating points, uncertainty, and synthetic factors.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 3.4 defines 500 paired full-size record-level resamples with frozen scores and observed budgets. Appendix B reports the 2.5th and 97.5th percentile ranges, while Figure 3 reports repetition frequencies for the controlled experiment. These quantities describe conditional sampling stability rather than retraining uncertainty.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: Model and simulation configurations and sample sizes are reported. Complete hardware specifications and end-to-end wall-clock totals for the historical model runs are not available in the archived records.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The study uses public-use deidentified vital-statistics records, releases no individual records or deployable clinical model, and frames results as evaluation evidence.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 5 discusses transparent resource accounting as a benefit and hidden workload expansion and redirected scarce services as operational risks.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work releases neither a deployable clinical model nor a new sensitive dataset.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Section 3 and Appendix D cite the NCHS Data User Agreement and the LightGBM MIT License. The package redistributes only derived aggregate results.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The anonymous supplement documents the audit inputs and outputs, exact commands, environment, aggregate source-data schema, and controlled-experiment code. It releases no individual-level dataset or pretrained clinical model.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study includes no crowdsourcing or direct participant interaction.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

675 Answer: [N/A]

676 Justification: The analysis uses deidentified public-use records and introduces no partici-

677 pant recruitment or intervention.

678 Guidelines:

- 679 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
- 680 with human subjects.
- 681 • Depending on the country in which research is conducted, IRB approval (or equiva-
- 682 lent) may be required for any human subjects research. If you obtained IRB approval,
- 683 you should clearly state this in the paper.
- 684 • We recognize that the procedures for this may vary significantly between institutions
- 685 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
- 686 guidelines for their institution.
- 687 • For initial submissions, do not include any information that would break anonymity
- 688 (if applicable), such as the institution conducting the review.

689 **16. Declaration of LLM usage**

690 Question: Does the paper describe the usage of LLMs if it is an important, original, or

691 non-standard component of the core methods in this research? Note that if the LLM is used

692 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

693 scientific rigor, or originality of the research, declaration is not required.

694 Answer: [N/A]

695 Justification: LLMs are not an experimental or methodological component of the study.

696 Guidelines:

- 697 • The answer [N/A] means that the core method development in this research does not
- 698 involve LLMs as any important, original, or non-standard components.
- 699 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
- 700 not be described.